

Problem Statement

 Update: candidate_pairs.tsv is part of your final submission

Blocking has to scale. Amazon resolves business entities across billions of records, so comparing every record with every other one is not an option. Your blocking / candidate-generation step must cut the search space to a small candidate set per Source 1 entity.

Candidate generation counts toward the final ranking. We will review your candidate_pairs.tsv and the code that produces it when deciding final rankings, alongside your matching_results.tsv score. The approach that generates a smaller candidate set per Source 1 entity will be ranked higher in the final evaluation beyond the public/private leaderboard.

Business Entity Resolution Challenge

In large-scale commercial platforms, business identity data arrives from multiple independent sources — each contributing partial, noisy fragments of information about the same real-world entities. These fragments share no common identifiers, and the challenge of determining which records refer to the same business is known as Entity Resolution (ER). Your challenge is to build an ML solution that, given business records from 3 independent data sources with noisy and inconsistent fields, determines which records across sources refer to the same real-world business entity.

Source 1 is the deduplicated reference source. Your task is to find all matching records from Source 2 and Source 3 for each Source 1 entity. A Source 1 entity may match zero, one, or many records from Source 2 and Source 3.

Download Data Set: [Click here to Download](#)

File Format

All files in this challenge are tab-separated (.tsv), and your submissions must be tab-separated too. Tabs are used because business addresses and the ID list columns both contain commas. Read them with an explicit tab separator, for example:

```
import pandas as pd
```

```
df = pd.read_csv("dataset/train/train_source1.tsv", sep="\t")
```

Reading a .tsv without sep="\t" will silently produce a single column containing the whole line.

Data Description:

Each source file (*_source1.tsv, *_source2.tsv, *_source3.tsv) has the following columns:

entity_id: Unique identifier for the record. The prefix indicates the source — S1-, S2-, or S3-.

business_name: Name of the business entity (may contain abbreviations, legal suffixes, typos, transliterations)

business_address: Address of the business (may contain partial addresses, format variations, missing components, landmark-based references)

country: Country label for the record. The training data covers US and India. The test set additionally contains a third country, France, that does not appear in the training data. Treat country as an open set of string labels: do not hard-code, filter, or one-hot your pipeline to only {US, India}, and remember that every test entity — France included — must appear in your submission.

There is no separate source column — a record's source is given by its entity_id prefix (S1-/S2-/S3-) and by which file it appears in.

The ground truth file (train_ground_truth.tsv) has two columns:

source1_entity_id: The entity_id of a Source 1 record

matched_entity_ids: Comma-separated list of matching entity_ids from Source 2 and/or Source 3 (empty when the entity has no matches)

Noise Patterns to Expect:

Name variations: Abbreviations (Corp vs. Corporation, Pvt vs. Private, Ltd vs. Limited), legal suffix inconsistencies, DBA/trade names, punctuation differences (& vs. "and"), word-order transpositions, typos

Address variations: Abbreviations (Rd vs. Road, St vs. Street), transliteration variants, missing components (no PIN code, no state), landmark-based references (Near SBI ATM), municipal numbering formats, component reordering

Dataset Details:

Training Dataset: Business records across 3 sources with ground truth matching labels

Test Set: Business records across 3 sources without matching labels

File Descriptions:

Training files

dataset/train/train_source1.tsv: Source 1 training records (the deduplicated reference source)

dataset/train/train_source2.tsv: Source 2 training records

dataset/train/train_source3.tsv: Source 3 training records

dataset/train/train_ground_truth.tsv: Ground truth matching labels for the training set

Test files

dataset/test/test_source1.tsv: Source 1 test records. Generate matches for every entity in this file.

dataset/test/test_source2.tsv: Source 2 test records

dataset/test/test_source3.tsv: Source 3 test records

No ground truth is provided for the test set. To measure your own performance, hold out a validation split from the training data and score it yourself using the F_{0.5} formula given below.

Output Format:

Your solution produces two tab-separated files, both placed in the output/ folder of your final submission package (see Final Submission Package below):

matching_results.tsv — your final entity matches. This is the only file scored on the leaderboard — it is what you upload to the Portal during the challenge.

candidate_pairs.tsv — the candidate set your blocking / candidate-generation stage produced, before your final matching model narrowed it down.

matching_results.tsv

Your final entity matches:

ColumnDescriptionsource1_entity_idThe entity_id of a Source 1

recordmatched_entity_idsComma-separated list of matching entity_ids from Source 2 and/or Source 3

Example (columns separated by a single tab, ID lists separated by commas with no quoting):

source1_entity_id	matched_entity_ids
S1-00001	S2-00047,S2-00193,S3-00812
S1-00002	S3-00004
S1-00003	

Important:

Every Source 1 entity in the test set must have exactly one row

Leave matched_entity_ids empty for entities with no matches (singletons)

No duplicate entity IDs within a single ID list

ID lists must only contain Source 2 or Source 3 IDs that exist in the test set

candidate_pairs.tsv

The candidate set from your blocking stage — every Source 2 / Source 3 record you considered a plausible match for each Source 1 entity, before your final matching model narrowed it down. This is the exact set of records you feed into your matching model for inference — the final candidate list just before the ML model scores them, not the raw output of an early blocking pass you later filter further. If your pipeline has several blocking/filtering stages, candidate_pairs.tsv is the last one: whatever your model actually runs inference over. Every ID in matching_results.tsv should therefore appear here.

It is not scored on the leaderboard; we use it to analyse blocking quality (recall ceiling, reduction ratio) and to verify your pipeline.

ColumnDescriptionsource1_entity_idThe entity_id of a Source 1

recordcandidate_entity_idsComma-separated list of candidate entity_ids from Source 2 and/or Source 3

Example:

source1_entity_id	candidate_entity_ids
S1-00001	S2-00047,S2-00193,S3-00812,S3-00999
S1-00002	S3-00004
S1-00003	

Same rules as matching_results.tsv: one row per Source 1 entity, candidate_entity_ids empty when blocking found no candidates, S2-/S3- IDs only, no duplicates within a list. Your final matches should be a subset of your candidates (a matched ID that never appeared as a candidate signals a pipeline bug — the validator warns about it).

Validate before submitting: a helper script utils/validate_submission.py (stdlib only, no dependencies) checks both files against every rule above so you can catch a rejection locally instead of spending a submission on it. Run it from this student_resource/ directory:

```
python3 utils/validate_submission.py \
  --matching output/matching_results.tsv \
  --candidate output/candidate_pairs.tsv \
  --test-dir dataset/test
```

It prints PASS (exit 0) when the files are safe to submit, or a numbered list of issues to fix (exit 1). It only reads your output files and the test source files; it does not compute your score.

Final Submission Package:

In addition to your live leaderboard uploads, every team submits a single zip archive with your code and outputs. We use it to reproduce your results, audit your blocking, and check the fair-play and model-license rules — the top teams' packages are reviewed in detail before the final rankings are confirmed.

Structure:

```
<team_name>_submission.zip
├── output/
│   ├── matching_results.tsv    # final matches (same file you upload to the leaderboard)
│   └── candidate_pairs.tsv     # your blocking candidate set
├── code/
│   ├── business_entity_resolution/
│   │   ├── src/               # all your source code
│   │   └── README.md          # how to reproduce end-to-end (data → blocking →
│                               matching → output)
```

| └─ requirements.txt # pinned dependencies / environment
| └─ Documentation_template.md # your methodology write-up (this filled-in template)
output/ — the two TSV files described above: matching_results.tsv and candidate_pairs.tsv.
code/business_entity_resolution/ — a self-contained, runnable copy of your pipeline. Put all
source under src/, and include a README.md with exact run instructions plus a
requirements.txt (or equivalent environment file) pinning versions. Anyone should be able to
regenerate both output files from the training/test data using only what is in this folder.
Methodology document — fill in the provided Documentation_template.md and drop it
straight into the zip (the filled-in .md is fine; a .pdf export works too). No need to rename it.
Constraints:

Format your output exactly as described above. Submissions that fail validation will not be
evaluated. You should see a SCORED status with your F_{0.5} score if the output is correctly
formatted.

matched_entity_ids must only reference entities from Source 2 or Source 3. Self-matches to
Source 1, and IDs that do not exist in the test set, will be rejected.

Every Source 1 entity must appear in your submission. Missing entities will cause rejection.
Duplicate entity IDs in any ID list will cause rejection, as will duplicate source1_entity_id
rows.

Final model should be a MIT/Apache 2.0 License model and up to 8 Billion parameters.

Evaluation Criteria:

Submissions are evaluated using F_β Score (β = 0.5) — a precision-heavy metric that
penalizes false merges (matching two different businesses) more than missed matches.

Formula:

$$F_{0.5} = (1.25 \times \text{Precision} \times \text{Recall}) / (0.25 \times \text{Precision} + \text{Recall})$$

Computed as a macro-average: F_{0.5} is calculated per Source 1 entity, then averaged
across all Source 1 entities in the evaluation set.

Singletons are included in that average. A Source 1 entity with no true matches scores 1.0
when you correctly predict an empty list, and 0.0 when you predict any match for it. Correctly
identifying singletons therefore earns credit, and false merges on them are penalised.

Why precision-heavy? In real-world entity resolution, merging two distinct businesses (false
positive) is more damaging than missing a link (false negative). F_{0.5} weights precision 2×
over recall.

Example:

Your model predicts S1-00001 matches [S2-00047, S2-00193, S3-00812]

Ground truth says S1-00001 matches [S2-00047, S3-00812]

Precision = 2/3, Recall = 2/2 = 1.0

$$F_{0.5} = (1.25 \times 0.667 \times 1.0) / (0.25 \times 0.667 + 1.0) = 0.714$$

Leaderboard Information:

Public Leaderboard: During the challenge, rankings will be based on a subset of the test set
to provide real-time feedback on your model's performance.

Private Leaderboard: After the challenge ends, the private leaderboard will be revealed,
which uses the remaining portion of the test set for evaluation.

Final Rankings: The final decision will be based on the private leaderboard.

You submit predictions for the full test set in both cases; the split is applied during scoring.

Submission Requirements:

Leaderboard (during the challenge): upload matching_results.tsv in the Portal —
tab-separated, with the exact column names described above. This is what drives the public
and private leaderboards.

Final submission package: submit the single zip described in Final Submission Package above — output/ with both matching_results.tsv (final matches) and candidate_pairs.tsv (your candidate-generation / blocking set fed to the model), code/business_entity_resolution/ (runnable pipeline), and your methodology document. All teams must submit it; the top teams' packages are reviewed before the final rankings are confirmed.

Your methodology document must describe:

- Methodology used

- Candidate generation/blocking strategy

- Model architecture and feature engineering

- Any other relevant information about the approach

A template for this documentation is provided in Documentation_template.md. There is no page limit — prioritise clarity and technical depth over brevity.

Academic Integrity and Fair Play:

 **STRICTLY PROHIBITED: External Data Lookup**

Participants are **STRICTLY NOT ALLOWED** to use external databases, APIs, or services to look up business identities or resolve entities. This includes but is not limited to:

- Using commercial entity resolution APIs or services

- Looking up business registrations from government databases

- Using geocoding APIs to normalize addresses

- Any external data augmentation from internet sources

Enforcement:

All submitted approaches, methodologies, and code pipelines will be thoroughly reviewed and verified

Any evidence of external data lookup will result in immediate disqualification

Fair Play: This challenge is designed to test your machine learning and data science skills using only the provided training data.

Tips for Success:

Invest in a strong blocking/candidate generation strategy — it determines the upper bound of your recall

Explore string similarity features (Jaccard, Levenshtein, TF-IDF cosine) for name and address matching

Pay attention to country-specific address patterns

Consider the precision-recall trade-off carefully — F_{0.5} rewards precision more than recall

Do not neglect singletons — correctly predicting "no match" is worth a full 1.0 on that entity

Validate your own output format against the rules above before submitting